

Evaluating Human Learning from RAG Systems

A literature orientation and design guide for a controlled study of source transparency, reliance, and knowledge acquisition in retrieval-augmented question answering

This guide organizes the literature your study will be evaluated against. It is arranged in six clusters, each of which corresponds to a claim you will have to defend in review. Every citation is annotated with what it gives you: a method to borrow, a result to predict against, a confound it warns about, or a gap it leaves open. Sections 7 through 10 turn the literature into a concrete design checklist, a measurement plan, and a validity-threat table.

The study, in one sentence

Participants learn an invented proprietary system (a fictional CLI or internal tool) through a RAG assistant over a synthetic corpus, under conditions that vary whether retrieved passages are visible and inspectable, and are then assessed on unassisted knowledge, transfer, and confidence calibration.

1. What you are actually claiming

Three separate contributions are in play, and reviewers will score them separately. Being explicit about which is primary determines which literature you must be exhaustive in and which you can merely gesture at.

Contribution	Primary claim	Cluster you must own
Empirical study	Source transparency changes what people learn, not just what they trust.	A, B, C
Methodological resource	A synthetic proprietary corpus enables contamination-free study of learning from RAG.	E
Design/system	A specific retrieval-time interaction improves outcomes.	D, F

The strongest single-paper framing is the empirical study, with the corpus as a methodological contribution reported alongside it and the interface manipulation kept minimal. A novel RAG architecture as a third contribution introduces a confound: if conditions differ in retrieval pipeline as well as interface, no learning difference can be attributed to transparency. Hold the backend fixed, hold the retrieved passages and generated answers identical across arms, and vary only what the participant can see and do.

2. Cluster A - Does AI assistance help or harm learning?

This is the literature that makes your research question legible. It also predicts the shape of your result, which means you should design to detect that shape rather than be surprised by it.

Bastani, H., Bastani, O., Sungu, A., Ge, H., Kabakci, O., & Mariman, R. (2025). Generative AI without guardrails can harm learning: Evidence from high school mathematics. *PNAS*, 122(26), e2422633122.

The single most important prior for your design. A randomized field experiment with roughly 1,000 Turkish high-school students across three arms: an unrestricted GPT-4 interface, a tutor-style interface with pedagogical safeguards, and a no-AI control. Assisted practice improved markedly in both AI arms; on a subsequent unassisted exam the unrestricted arm scored about 17% below control, an effect the guardrailed arm eliminated. Two design lessons transfer directly: (1) measure assisted performance and unassisted performance separately, because they dissociate and can point in opposite directions; (2) interface design, not model capability, is what moved the learning outcome - which is exactly the logic of your transparency manipulation. Note the 2025 correction notice when citing.

Kazemitabaar, M., Chow, J., Ma, C. K. T., Ericson, B. J., Weintrop, D., & Grossman, T. (2023). Studying the effect of AI code generators on supporting novice learners in introductory programming. *CHI '23*.

The methodological template for the learning half of your study, and a useful counterweight to Bastani. 69 novices, 45 authoring tasks paired with unassisted modification tasks, an immediate post-test and a retention test one week later. Codex access raised authoring performance without degrading manual modification; the delayed retention advantage was positive but not significant. Borrow the paired assisted/unassisted task structure and the one-week delayed test. The divergence from Bastani is itself worth a paragraph in your related work: the effect of AI assistance on learning is contested, which is what makes another well-controlled study valuable.

Lee, H.-P., Sarkar, A., Tankelevitch, L., Drosos, I., Rintel, S., Banks, R., & Wilson, N. (2025). The impact of generative AI on critical thinking: Self-reported reductions in cognitive effort and confidence effects from a survey of knowledge workers. *CHI '25*.

319 knowledge workers, 936 first-hand task examples. The key quantitative finding is a double dissociation of confidence: higher confidence in the AI predicts less enacted critical thinking, while higher self-confidence in the task domain predicts more. This is your justification for treating confidence as a measured variable rather than a nuisance, and for collecting a prior-domain-experience covariate. It also grounds the enterprise onboarding framing in real workplace data rather than a hypothetical.

Zhai, C., Wibowo, S., & Li, L. D. (2024). The effects of over-reliance on AI dialogue systems on students' cognitive abilities: A systematic review. *Smart Learning Environments*, 11(1), 28.

Use as a one-citation summary of the overreliance-in-education literature so your related work section does not have to enumerate it. Helpful for establishing that the concern is systematic rather than anecdotal.

Design implication

Report at least two dependent variables that can dissociate: performance *with* the system and knowledge *without* it. A study that only measures assisted task success cannot distinguish a good tool from a good crutch, and reviewers who know the Bastani result will say so.

3. Cluster B - Reliance, transparency, and verification

This cluster establishes the non-obviousness of your hypothesis. The naive prediction is that showing sources helps. The literature says transparency frequently raises trust without raising accuracy, and sometimes raises acceptance of wrong answers. That tension is your paper's engine.

Bucinca, Z., Malaya, M. B., & Gajos, K. Z. (2021). To trust or to think: Cognitive forcing functions can reduce overreliance on AI in AI-assisted decision-making. *PACM HCI*, 5(CSCW1), Article 188.

The canonical demonstration that interventions requiring effortful engagement outperform simply presenting explanations, and that participants trust and prefer the interfaces they perform worse with. Two things to carry forward. First, the subjective/objective divergence: collect preference and trust ratings, expect them to favor the condition that produces weaker learning, and treat that divergence as a finding. Second, Need for Cognition moderated the benefit, which is a strong argument for including an NFC scale as a covariate and auditing for differential benefit across participants.

Liu, N. F., Zhang, T., & Liang, P. (2023). Evaluating verifiability in generative search engines. *Findings of EMNLP 2023*.

The audit establishing that citations in generative search frequently fail to support the statements they are attached to - roughly half of generated statements were fully supported by their citations in their sample. This is the empirical basis for your unsupported-claim probe items: if systems in the wild produce plausible answers that their own sources do not support, then whether users can detect that is a real and consequential ability to measure.

Fostering appropriate reliance on large language models: The role of explanations, sources, and inconsistencies. (2025). *CHI '25*. DOI 10.1145/3706598.3714020.

The closest existing work to your transparency manipulation, and the paper a reviewer is most likely to ask how you differ from. A pre-registered experiment (N=308) using pre-constructed responses so that response features could be fully controlled, finding that participants seek out sources to verify details and rely less on incorrect answers when sources are informative. Your differentiation: they measure single-shot reliance on factual questions; you measure cumulative knowledge acquisition over a learning session in a domain with no parametric prior. Read their stimulus-construction method closely - controlling responses in advance is the right move and you should consider adopting it.

To rely or not to rely? Evaluating interventions for appropriate reliance on large language models. (2025). *CHI '25*. DOI 10.1145/3706598.3714097.

Pre-registered, N=400, three reliance interventions across two task types. The headline is sobering and useful to you: interventions reduced over-reliance but generally failed to improve *appropriate* reliance. Cite this to pre-empt the reviewer question of why anyone should expect your manipulation to work, and to justify measuring appropriate reliance with a metric that separates it from raw agreement.

Effects of LLM-based search on decision making: Speed, accuracy, and overreliance. (2025). *CHI '25*. DOI 10.1145/3706598.3714082.

Establishes the speed/quality trade-off in LLM-mediated search and reports participants' own articulation that absent citations made reliability hard to assess. Useful for motivating the non-RAG search baseline arm.

The measure that distinguishes your paper

Availability is not consumption. Log expansion and dwell on retrieved passages separately from whether the passages were shown. If the transparency arm learns more only among participants who actually opened sources, the finding is about inspection; if it holds regardless, the finding is about exposure. Most papers in this cluster cannot make that distinction. You can.

4. Cluster C - Metacognition, calibration, and the illusion of knowing

This is where the most interesting version of your result lives, and it is the cluster most likely to be missing from a systems-oriented related-work section. A fluent RAG answer is an efficient machine for producing unearned feelings of understanding.

Fisher, M., Goddu, M. K., & Keil, F. C. (2015). Searching for explanations: How the Internet inflates estimates of internal knowledge. *Journal of Experimental Psychology: General*, 144(3), 674-687.

Nine experiments showing that searching online for explanations raises people's estimates of their own internal knowledge - they conflate accessible information with personally held understanding, and the inflation persists into unrelated domains. The effect survived controls for time, content, and search autonomy, and appeared even when the search returned nothing useful. This is the theoretical basis for predicting a confidence-accuracy gap in your opaque condition, and arguably in both conditions relative to a read-the-documents control.

Rozenblit, L., & Keil, F. (2002). The misunderstood limits of folk science: An illusion of explanatory depth. *Cognitive Science*, 26(5), 521-562.

The parent effect. People believe they understand mechanisms in far more detail than they do, and the illusion collapses when they attempt to produce an explanation. Directly motivates including a free-response explanation item in your post-test rather than only multiple choice: the explanation-generation task is the instrument that exposes the gap.

Bjork, R. A., Dunlosky, J., & Kornell, N. (2013). Self-regulated learning: Beliefs, techniques, and illusions. *Annual Review of Psychology*, 64, 417-444.

The authoritative review of why learners misjudge their own learning, with fluency as the central culprit: ease of processing in the moment is read as evidence of durable learning, and it is a bad signal. Everything RAG does well - immediacy, synthesis, coherence - increases processing fluency. This paper lets you state the mechanism your study tests in one sentence.

Bjork, E. L., & Bjork, R. A. (2011). Making things hard on yourself, but in a good way: Creating desirable difficulties to enhance learning. In *Psychology and the Real World*, 56-64.

The framing device for arguing that source inspection might function as a desirable difficulty: a manipulation that slows the interaction and reduces immediate satisfaction while improving retention and transfer. If your transparency arm is slower, less liked, and better retained, this is the theory that makes the pattern coherent rather than anomalous.

Roediger, H. L., & Karpicke, J. D. (2006). Test-enhanced learning: Taking memory tests improves long-term retention. *Psychological Science*, 17(3), 249-255.

The testing effect, and the reason your immediate post-test contaminates your delayed post-test if you are not careful. Also the reason immediate testing systematically overstates learning: in the original demonstration the re-study advantage at five minutes reversed by two days and reversed sharply by one week. If you can only afford one measurement point, the delayed one is more informative; if you take both, do not treat them as interchangeable.

Recommended headline measure

Per-item confidence collected alongside each post-test answer, analyzed as calibration (for example a Brier score or over/underconfidence bias score) rather than as mean confidence. A result showing equal accuracy but worse calibration in the opaque condition is more publishable, and more consequential for enterprise onboarding, than a small main effect on quiz score. An overconfident junior engineer is worse than an uncertain one.

5. Cluster D - Interaction and information-seeking foundations

These give your interface manipulation a theoretical lineage rather than making it look like an arbitrary UI variant. They are also what makes the work legible as HCI rather than as an NLP evaluation with humans attached.

Horvitz, E. (1999). Principles of mixed-initiative user interfaces. *CHI '99*, 159-166.

The foundational statement of when a system should act autonomously versus hand control back, including using dialogue to resolve key uncertainties and weighing the cost of interrupting. If you pursue the ambiguity-triggered variant - intervening at retrieval only when the candidate set is genuinely ambiguous - this is the framework, and framing it as applying Horvitz's arbitration principles to a substrate that did not exist in 1999 is a clean contribution statement.

Amershi, S., Weld, D., Vorvoreanu, M., Fourney, A., Nushi, B., Collisson, P., Suh, J., Iqbal, S., Bennett, P. N., Inkpen, K., Teevan, J., Kikin-Gil, R., & Horvitz, E. (2019). Guidelines for human-AI interaction. *CHI '19*.

The consolidated design guidelines, several of which bear directly on your conditions (make clear why the system did what it did; support efficient correction; scope services to context). Useful for deriving your interface variants from stated principles instead of intuition, which strengthens the design-rationale section.

Bates, M. J. (1989). The design of browsing and berrypicking techniques for the online search interface. *Online Review*, 13(5), 407-424.

The berrypicking model: real information seeking is an evolving process in which the need is refined by encountering results, gathered bit by bit rather than satisfied by a single matched set. RAG collapses this into one synthesized answer and discards the candidate set. This is the sharpest theoretical statement of what a non-transparent RAG interface takes away from a learner, and it predates the technology by 35 years, which makes it rhetorically effective.

Marchionini, G. (2006). Exploratory search: From finding to understanding. *Communications of the ACM*, 49(4), 41-46.

Distinguishes lookup from learning and investigation as search activities with different interface requirements. Your task is explicitly a learning search, not a lookup, and this is the citation that says those are not the same thing and should not be evaluated the same way.

White, R. W., & Roth, R. A. (2009). *Exploratory search: Beyond the query-response paradigm*. Morgan & Claypool.

Book-length treatment including evaluation methodology for exploratory search systems, where success is not a single relevance judgment. Useful when you defend using learning outcomes rather than answer accuracy as the dependent variable.

Belkin, N. J., Oddy, R. N., & Brooks, H. M. (1982). ASK for information retrieval: Part I. *Journal of Documentation*, 38(2), 61-71.

Anomalous states of knowledge: users cannot specify what they need precisely, because the inability to specify it is the problem. This is the theoretical justification for the ambiguity-arbitration idea - queries are underspecified by nature, and RAG silently resolves that underspecification without telling anyone.

6. Cluster E - Contamination-resistant synthetic corpora

This is your methodological contribution and the least crowded cluster. Almost nobody studying AI-assisted learning can guarantee zero prior knowledge and zero model-parametric contamination; a well-built invented domain gives you both, plus complete ground truth over what is and is not answerable.

RIKER: Scalable and reliable evaluation of AI knowledge retrieval systems via a coherent simulated universe. arXiv:2601.08847.

The closest thing to a template for what you want to build. A ground-truth-first methodology: define structured ground truth, then generate synthetic documents from it, inverting the usual pipeline of annotating existing documents. This yields contamination resistance and exact answerability labels by construction. Their finding that grounding ability and hallucination resistance are distinct capabilities - a system can be good at finding facts that exist while still fabricating facts that do not - is a direct argument for including unanswerable items in your quiz.

LongBioBench: A controllable examination for long-context language models. arXiv:2506.02921.

Uses fictional biographies as a controlled playground, and articulates the design criteria you should adopt and cite: target information must be embedded fluently and coherently rather than inserted as a detectable needle; the benchmark must be configurable; and the task must be free of parametric knowledge to prevent contamination. Steal this criteria list for your corpus-design section.

AntiLeakBench: Preventing data contamination by automatically constructing benchmarks with updated real-world knowledge. arXiv:2412.13670.

Argues that merely-recent data does not guarantee contamination-free evaluation, since new collections still contain pre-existing knowledge, and that explicitly constructed novel knowledge is required. Use this to justify invention over obscurity: picking a genuinely obscure real tool is not equivalent to inventing one.

Fictitious synthetic data can improve LLM factuality via prerequisite learning. arXiv:2410.19290.

Method for generating fictitious entities with verified zero overlap against reference corpora, including the filtering step that checks generated names against existing entries. Practically useful: this is the naming-collision check you need to run before committing to your invented tool's vocabulary.

From raw corpora to domain benchmarks: Automated evaluation of LLM domain expertise. arXiv:2506.07658.

A deterministic corpus-to-benchmark pipeline that avoids LLM judges and manual curation. Relevant if you want your quiz-generation procedure to be reproducible and auditable rather than hand-authored, which strengthens the resource contribution considerably.

Corpus construction requirements, drawn from this cluster

1. Invented vocabulary with verified zero collision against real tools, checked by adversarially probing the model with your invented names before the study.
2. Realistic pathology: a deprecated document contradicting a current one, an internal thread containing a confidently wrong answer, a partially documented flag. A clean corpus tests nothing interesting about source inspection.
3. Explicit answerability labels per item, including items with no answer in the corpus, so you can measure detection of unsupported answers rather than only recall.
4. Independent validation: someone who did not build the corpus confirms every quiz item is answerable from it.
5. Release the corpus, the generation procedure, and the item blueprint - that is what makes it a resource rather than an appendix.

7. Cluster F - Human-centered RAG (the immediate neighbors)

A short cluster, deliberately. These are the papers that establish RAG-specific HCI as an active area and that a reviewer will expect you to have read; none of them measures learning, which is your gap.

Exploring trust and transparency in retrieval-augmented generation for domain experts. (2025). *CHI '25 Extended Abstracts*. DOI 10.1145/3706599.3719985.

User-centered RAG design for financial domain experts, with an empirical evaluation of transparency-enhancing features against trust and understanding. The nearest neighbor to your framing. Note what it measures - trust and perceived understanding - and note that measured knowledge is exactly what it does not measure. That sentence belongs in your introduction.

RAGTrace: Understanding and refining retrieval-generation dynamics in retrieval-augmented generation. arXiv:2508.06056.

Visual-analytics approach to making the retrieval-to-generation relationship inspectable. Useful as a source of interface primitives for your transparency condition, and as evidence that retrieval-stage intervention is a live design question rather than a solved one.

Hoque, M. N., Mashiat, T., Ghai, B., Shelton, C. D., Chevalier, F., Kraus, K., & Elmqvist, N. (2024). The HaLLMark effect: Supporting provenance and transparent use of large language models in writing with interactive visualization. *CHI '24*.

Provenance surfacing in a generative writing tool. Adjacent domain, transferable interface vocabulary, and a good citation for the general claim that provenance transparency is an established HCI concern rather than an idiosyncratic one.

8. Design checklist derived from the literature

Each item below traces to a specific source above. If you can satisfy all of them, the study is defensible; the ones you cannot satisfy should appear in your limitations section by name.

- **One backend, minimal interface difference.** Identical retrieved passages and identical generated answers across arms; vary only visibility and inspectability. Anything else confounds interface with pipeline.
- **At least three arms.** Opaque RAG, transparent RAG, and a non-RAG baseline (plain search or the documents alone) over the same corpus. Without the baseline, the transparency comparison has no reference point and reviewers will ask whether RAG beats reading. (Bastani et al. use exactly this no-AI control.)
- **Paired assisted and unassisted measurement.** An assisted task during the session and a closed-book test after, following Kazemitabaar et al. and Bastani et al. Report both; expect them to dissociate.
- **A delayed post-test at one week** if attrition budget allows. Roediger and Karpicke's reversal is the reason immediate-only measurement overstates learning.
- **Pre-constructed responses.** Consider fixing the assistant's answers in advance, as the CHI 2025 reliance work did, so that response quality is controlled rather than sampled.
- **Fixed-time rather than fixed-task sessions,** to blunt the exposure confound: the transparency arm reads more text and spends more time, and time is the trivial alternative explanation for any advantage.
- **Interaction logging that separates availability from consumption** - expansions, dwell, revisits, per source.
- **Unanswerable and contradicted items in the quiz,** so you can measure detection of unsupported answers (motivated by Liu et al. and RIKER).

- **Per-item confidence for calibration**, plus at least one free-response explanation item (Rozenblit and Keil).
- **Covariates**: prior programming or domain experience, and Need for Cognition (Bucinka et al. found NFC moderated intervention benefit; auditing for differential benefit is now expected).
- **Pre-registration** of hypotheses, primary DV, exclusion rules, and analysis plan. Both closely comparable CHI 2025 papers were pre-registered; an unregistered comparison will look weaker by contrast.
- **Power planning**. Between-subjects learning outcomes are noisy; plan roughly 40-50 per cell for a medium effect and pilot for quiz difficulty, since ceiling and floor effects silently destroy these designs.

9. Measurement plan

Construct	Instrument	Why, and source
Assisted performance	Task completion and correctness during the session, with the system available	Detects the crutch effect; the arm where AI always looks good (Bastani et al.)
Unassisted knowledge	Closed-book post-test: recall, application, cross-document synthesis items	The actual learning outcome; synthesis items are where source inspection should pay off
Transfer	Novel task requiring composition of learned concepts, not covered verbatim in corpus	Distinguishes memorized answers from usable knowledge (Bjork and Bjork)
Retention	Delayed post-test at one week	Immediate testing overstates fluent scaffolded learning (Roediger and Karpicke)
Calibration	Per-item confidence, scored as bias and Brier score	The predicted dissociation: equal accuracy, worse calibration (Fisher et al.)
Illusion of depth	Free-response mechanism explanation, rated by blinded coders	Exposes the gap that multiple choice hides (Rozenblit and Keil)
Unsupported-answer detection	Items whose answers are absent from or contradicted in the corpus	Verification ability, the thing transparency is supposed to enable (Liu et al.)
Source engagement	Expansion counts, dwell time, revisits, per retrieved passage	Separates availability from consumption; unique to your design
Subjective experience	Trust, perceived learning, preference, workload (e.g. NASA-TLX)	Expect divergence from objective outcomes; that divergence is a finding (Bucinka et al.)
Covariates	Prior domain experience; Need for Cognition	Dominates variance; moderates intervention benefit

10. Threats to validity, and what to do about each

Threat	Why it bites	Mitigation
Exposure confound	Transparency arm reads more and spends longer; time explains any advantage	Fixed-time sessions; log and covary reading time; add an arm where sources are shown but not expandable
Parametric leakage	Model recognizes your invented tool or blends in a real one	Collision-check names; adversarially probe the model pre-study; run a no-corpus control query set
Ceiling / floor	45-minute session on an invented CLI is easy to make trivial or impossible	Pilot for difficulty; blueprint items across recall, application, synthesis; pre-register exclusion rules
Testing effect contamination	Immediate post-test acts as retrieval practice and inflates the delayed test	Counterbalance item sets across measurement points; or accept and report it as a constant across arms
Demand characteristics	Participants in the transparency arm infer sources are the point and perform inspection	Neutral framing; do not name the manipulation; measure spontaneous engagement rather than instructed
Ecological validity	One session with a fictional tool is not enterprise onboarding	Say so plainly; the trade is internal validity for realism, and the contamination control is what you bought
Underpowered null	A null on learning is uninformative if the design could not have detected an effect	Power analysis in the pre-registration; report equivalence tests or Bayes factors for nulls

11. Where the contribution is, and where it is not

Framings that will be treated as weak: a novel RAG pipeline evaluated on retrieval metrics with a user study attached; a system paper whose evaluation is satisfaction ratings; a study whose only outcome measure is whether participants liked the interface. Each of these is a result that cannot embarrass you, which is another way of saying it cannot inform anyone.

Framings that hold up: an unassisted learning outcome that could have gone either way; a calibration result showing confidence and knowledge coming apart; a released corpus that other groups can use to run contamination-free studies; a demonstration that the availability of sources and the consumption of sources have different effects. Note that the third of these is a contribution even if the study's main effect is null, which is a useful property in a first project.

On venues: IUI fits the interaction-mechanism framing and has a March conference cycle; CHI fits the reliance and learning experiment; the CHI Late-Breaking Work track or a workshop is a reasonable first outing for the corpus methodology alone; CHIIR and SIGIR fit if you frame it as interactive information retrieval, which the Bates and Marchionini lineage supports. A resource or dataset track is worth considering for the corpus as a separable artifact.

Scoping advice

Two arms, one backend, a strong corpus, and calibration as the headline dependent variable is a complete and defensible paper. A novel system, four arms, a delayed post-test, and a new evaluation standard is a dissertation chapter. Decide which you are writing before you start building, and have that conversation with your advisor early - the answer determines how much of the corpus work is worth automating.

12. Consolidated reference list

Learning under AI assistance

- Bastani, H., Bastani, O., Sungu, A., Ge, H., Kabakci, O., & Mariman, R. (2025). Generative AI without guardrails can harm learning: Evidence from high school mathematics. *PNAS*, 122(26), e2422633122. <https://doi.org/10.1073/pnas.2422633122> (see also correction, *PNAS* 122(34), 2025)
- Kazemitabaar, M., Chow, J., Ma, C. K. T., Ericson, B. J., Weintrop, D., & Grossman, T. (2023). Studying the effect of AI code generators on supporting novice learners in introductory programming. *CHI '23*. <https://doi.org/10.1145/3544548.3580919>
- Lee, H.-P., Sarkar, A., Tankelevitch, L., Drosos, I., Rintel, S., Banks, R., & Wilson, N. (2025). The impact of generative AI on critical thinking. *CHI '25*. <https://doi.org/10.1145/3706598.3713778>
- Zhai, C., Wibowo, S., & Li, L. D. (2024). The effects of over-reliance on AI dialogue systems on students' cognitive abilities: A systematic review. *Smart Learning Environments*, 11(1), 28.

Reliance, transparency, verification

- Bucinka, Z., Malaya, M. B., & Gajos, K. Z. (2021). To trust or to think. *PACM HCI*, 5(CSCW1), Article 188. <https://doi.org/10.1145/3449287>
- Liu, N. F., Zhang, T., & Liang, P. (2023). Evaluating verifiability in generative search engines. *Findings of EMNLP 2023*. arXiv:2304.09848
- Fostering appropriate reliance on large language models: The role of explanations, sources, and inconsistencies. (2025). *CHI '25*. <https://doi.org/10.1145/3706598.3714020>
- To rely or not to rely? Evaluating interventions for appropriate reliance on large language models. (2025). *CHI '25*. <https://doi.org/10.1145/3706598.3714097>
- Effects of LLM-based search on decision making: Speed, accuracy, and overreliance. (2025). *CHI '25*. <https://doi.org/10.1145/3706598.3714082>

Metacognition and calibration

- Fisher, M., Goddu, M. K., & Keil, F. C. (2015). Searching for explanations: How the Internet inflates estimates of internal knowledge. *JEP: General*, 144(3), 674-687. <https://doi.org/10.1037/xge0000070>
- Rozenblit, L., & Keil, F. (2002). The misunderstood limits of folk science: An illusion of explanatory depth. *Cognitive Science*, 26(5), 521-562.
- Bjork, R. A., Dunlosky, J., & Kornell, N. (2013). Self-regulated learning: Beliefs, techniques, and illusions. *Annual Review of Psychology*, 64, 417-444.
- Bjork, E. L., & Bjork, R. A. (2011). Making things hard on yourself, but in a good way. In *Psychology and the Real World* (pp. 56-64). Worth Publishers.
- Roediger, H. L., & Karpicke, J. D. (2006). Test-enhanced learning. *Psychological Science*, 17(3), 249-255.

Interaction and information-seeking foundations

- Horvitz, E. (1999). Principles of mixed-initiative user interfaces. *CHI '99*, 159-166. <https://doi.org/10.1145/302979.303030>
- Amershi, S., et al. (2019). Guidelines for human-AI interaction. *CHI '19*. <https://doi.org/10.1145/3290605.3300233>
- Bates, M. J. (1989). The design of browsing and berrypicking techniques for the online search interface. *Online Review*, 13(5), 407-424. <https://doi.org/10.1108/eb024320>
- Marchionini, G. (2006). Exploratory search: From finding to understanding. *CACM*, 49(4), 41-46.

White, R. W., & Roth, R. A. (2009). *Exploratory search: Beyond the query-response paradigm*. Morgan & Claypool.

Belkin, N. J., Oddy, R. N., & Brooks, H. M. (1982). ASK for information retrieval: Part I. *Journal of Documentation*, 38(2), 61-71.

Synthetic corpora and contamination control

RIKER: Scalable and reliable evaluation of AI knowledge retrieval systems and the coherent simulated universe. arXiv:2601.08847

LongBioBench: A controllable examination for long-context language models. arXiv:2506.02921

AntiLeakBench: Preventing data contamination by automatically constructing benchmarks with updated real-world knowledge. arXiv:2412.13670

Fictitious synthetic data can improve LLM factuality via prerequisite learning. arXiv:2410.19290

From raw corpora to domain benchmarks: Automated evaluation of LLM domain expertise. arXiv:2506.07658

Human-centered RAG

Exploring trust and transparency in retrieval-augmented generation for domain experts. (2025). *CHI '25 EA*. <https://doi.org/10.1145/3706599.3719985>

RAGTrace: Understanding and refining retrieval-generation dynamics in retrieval-augmented generation. arXiv:2508.06056

Hoque, M. N., et al. (2024). The HaLLMark effect: Supporting provenance and transparent use of large language models in writing with interactive visualization. *CHI '24*. <https://doi.org/10.1145/3613904.3641895>

Verification note: citations were checked against publisher records, arXiv, or indexing services at the time of writing. Several 2025-2026 CHI and arXiv entries are listed without full author lists where the source record was partial; confirm author order and page numbers against the ACM Digital Library before submission. For the two most recent arXiv preprints, check whether a peer-reviewed version has since appeared and cite that instead.